



Caching and State Management

Google App Engine

Agenda

1

Edge Caching

2

Memcache Overview

3

Implementing Memcache

4

Caveats and Solutions




Edge Caching

Introduction

Edge Caching at Google Frontends

GFE closest to the user provides edge caching capability



Enable Edge Caching for App Engine Apps

Option 1 - Use cache-control header on HTTP Response

```
class MyHandler(webapp.RequestHandler):  
    def get(self):  
        self.response.headers.add_header(  
            'cache-control',  
            'public, max-age=7200') # 2hr cache
```

May be honored by Google Frontend
(not guaranteed)

Enable Edge Caching for App Engine Apps

Option 2 - Define Static Content

Java

appengine-web.xml

```
...  
<static>  
  <include path="/**/*.png" />  
  <exclude path="/img/**/*.png" />  
</static>  
...
```

Python

app.yaml

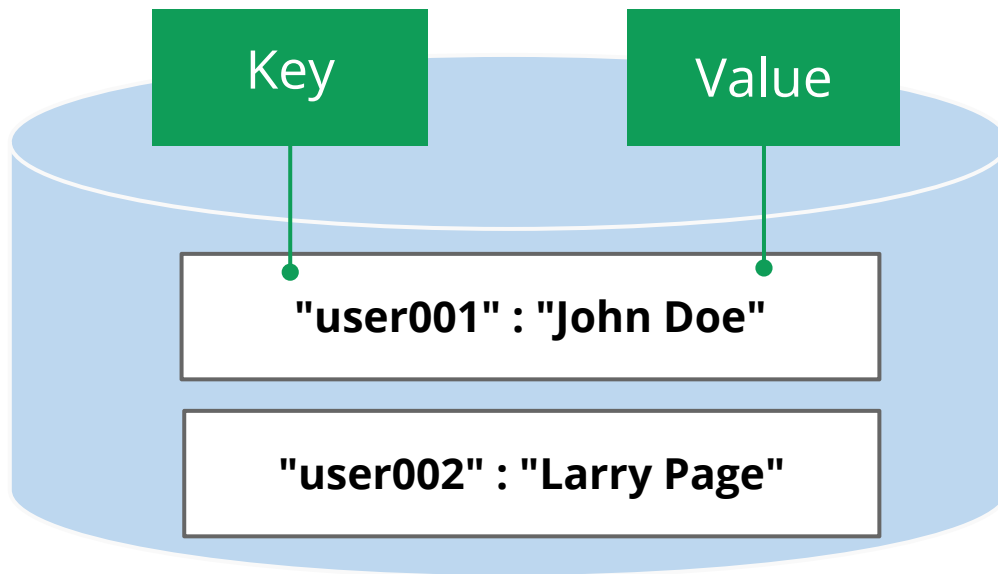
```
...  
- url: /images  
  static_dir: static/images  
OR  
- url: /images/(.*)  
  static_files: static/images/\1  
  upload: static/images/(.*)  
...
```



Memcache

Introduction

What Is Memcache?



Applications for Memcache

- Caching in front of Datastore
 - Cache entities for low-latency reads
 - Integrated into most ORM frameworks (ndb, Objectify)
- Caching for Read-heavy operations
 - User authentication token and session data
 - APIs call or other computation results
- Semi-durable shared-state cross app instances
 - Sessions
 - Counters/Metrics
 - Application Configurations



Implementing Memcache

Why Use Memcache?

Why Do We Use Memcache?

Time and Money!

1. Improve Application Performance
2. Reduce Application Cost

Memcache Speed

How fast can Memcache be?

Query: Latency

Measures the latency, in milliseconds, of one call to `datastore.query()` on an High Replication Datastore.



Datastore Query Latency

Get: Latency

Measures the latency, in milliseconds, of call to `memcache.get()` for a single item.



Memcache Read Latency

Memcache APIs

Java

- JCache APIs
- GAE Low-level Memcache APIs
- Objectify for Datastore

Python

- `google.appengine.api.memcache` module
- **NDB** for Datastore

General Memcache Usage Pattern

Coordinate data *read* with Datastore

- ❑ Check if Memcache value exists
 - If it does, display/use cached value directly; otherwise fetch the value from Datastore and write the value to Memcache

Coordinate data *write* with Datastore

- ❑ Update Memcache value
 - To handle race condition, leverage put if untouched/compare and set to detect race conditions
- ❑ Write the value to Datastore
 - Optionally, leverage the task queue for background writes

Using GAE Client Library

python

```
from google.appengine.api import memcache
...
value = memcache.get(key)

if value is None:
    value = get_value_from_db(key)
    if not memcache.add(key, value):
        logging.error('Memcache add failed.')
...
```

Batch Operations

getAll (), putAll(), deleteAll()

- A single read or write operation for multiple memcache entries
- Further improve Memcache performance by reducing network traffic
- Batch size < 32 MB

Dedicated Memcache

Memcache Service:

The memcache Shared class is free of charge, but provided on a "best effort" basis with no space guarantees. The Dedicated class assigns resources exclusively to your application and is billed by the gigabyte-hour of cache space (requires billing to be enabled). [Learn more](#).

Class **Dedicated (10k ops/s/GB)** of size GB (range 1-20 GB) , 50k ops/s aggregate

Note: Saving this change will cause a full flush of your cache.

1. Provides fixed cache capacity that is exclusive to your application.
2. Billed by the GB per hour of cache size. Charged in 15 minute increments.
3. Gives you additional control over cache size. More predictable performance

Note: Whether shared or dedicated, Memcache is not a durable storage. Plan for your application to function without Memcache.

Key Takeaways


1. Memcache is supported natively in GAE. Take advantage of it to improve your GAE application performance.
2. Memcache supports open standard JCache API. Many advanced features are available by GAE Memcache APIs to suit your application's need i.e. Batch, Atomic, Asynchronous operations
3. Seamless integration with GAE Datastore in a few libraries like Python ndb and Java Objectify

Key Takeaways

4. Read-frequently and write-rarely data is most suitable in combining with Memcache
5. Handle Memcache's volatility in your application
6. Use Memcache wisely, it is not an unlimited resource

Hands-on

- You will create an App Engine application in either Java or Python
- You will then integrate the Memcache API to understand how it works



Images by Connie Zhou

Resources

- Memcache Java API Overview
<https://cloud.google.com/appengine/docs/java/memcache/>
- Memcache Python API Overview
<https://cloud.google.com/appengine/docs/python/memcache/>



cloud.google.com