



# Datastore Queries and Indexes

Google App Engine



Google Cloud Platform

# Agenda

1

Queries

2

Indexes

3

Restrictions and Workarounds for Queries

4

Cost of Indexing



# Introduction to Queries

# Query

- Query by **kind**
- Filter on **property** or **parent**
- Results can be **sorted**

## Consistency Model

Queries always return results from the **INDEX**

- Eventually Consistent
  - Queries without an ancestor filter get results from the last index update
- Strongly Consistent
  - Queries using an ancestor filter forces applicable index updates to complete
  - `get()` forces the index to update before returning the results

# Filters

## Filter on

- property values
- keys
- ancestors

Use an Ancestor  
Filter for Strong  
Consistency

## Filter on property values

- Equality filter (Equal to)
- IN -- Member of a list
- Inequality filters
  - not equal to
  - less than
  - less than or equal to
  - greater than
  - greater than or equal to

# Query

- Adding multiple filters to a query

## Java

```
Query q = new Query("Person");  
Query.Filter filter1 = new FilterPredicate(...);  
Query.Filter filter2 = new FilterPredicate(...);  
Query.Filter comboFilter = CompositeFilterOperator.and(filter1, filter2);  
q.setFilter(comboFilter);
```

## Python

```
q = Person.all()  
q.filter("name =", "John")  
q.filter("sport =", "Sailing")
```

# Sort

- Query results can be sorted by property
  - Sort by **ascending** or **descending** value of a property
  - Some restrictions on sorting (discussed later)

## Java

```
q.addSort("name");
```

## Python

```
q.order('name')
```



# Indexes



# Index

Bigtable doesn't support  
Query. How can  
Datastore support it?

In Datastore, Queries are  
executed as *Index Scans* on  
Bigtable.

# Index

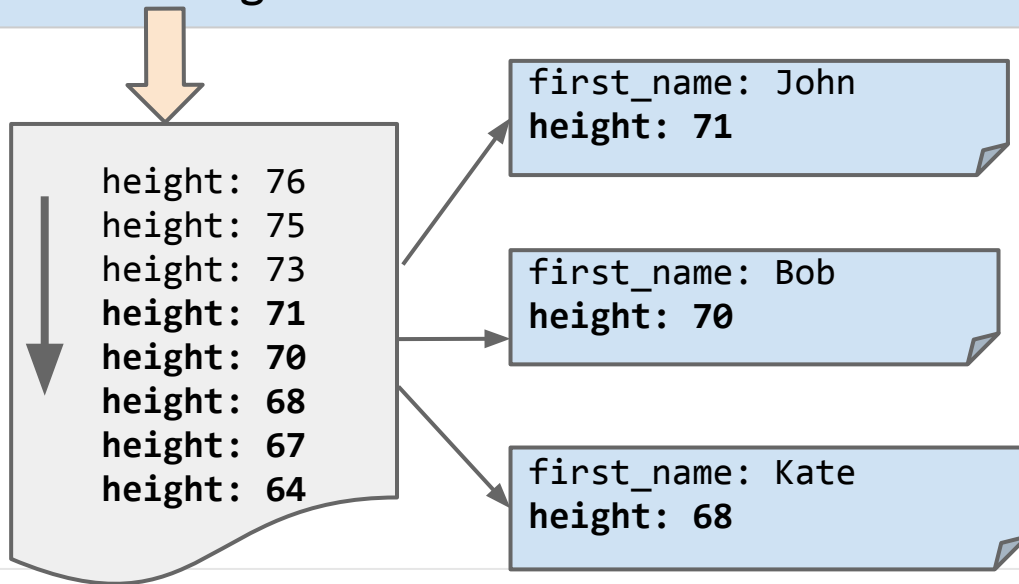
- Queries are executed as index scans

```
SELECT * FROM Person  
WHERE height < 72  
ORDER BY height DESC
```

**Datastore  
Query**

**Index table  
for height**

**Range Scan  
on Bigtable**



**Entities  
in the  
query  
result**

# Index

Datastore requires **index** for ALL queries, otherwise the **query fails**

↳ Not like the index in RDBMS, which is used to improve performance



The value of **Index Scan**

The **Index Scan** makes it possible for **query performance to scale with the size of the result set**, not the data set.

# Index

## Single-property index

- Query for kind = Person, first\_name = 'Ben'
- Scan with prefix [Person first\_name Ben]

Single-property  
index is created  
automatically



Kind	Property	Value	Entity Key
Person	first_name	Alice	EntityKey112
Person	first_name	Ben	EntityKey122
Person	first_name	Bridgit	EntityKey223
Person	first_name	Cathy	EntityKey233

Key of Index table

Value

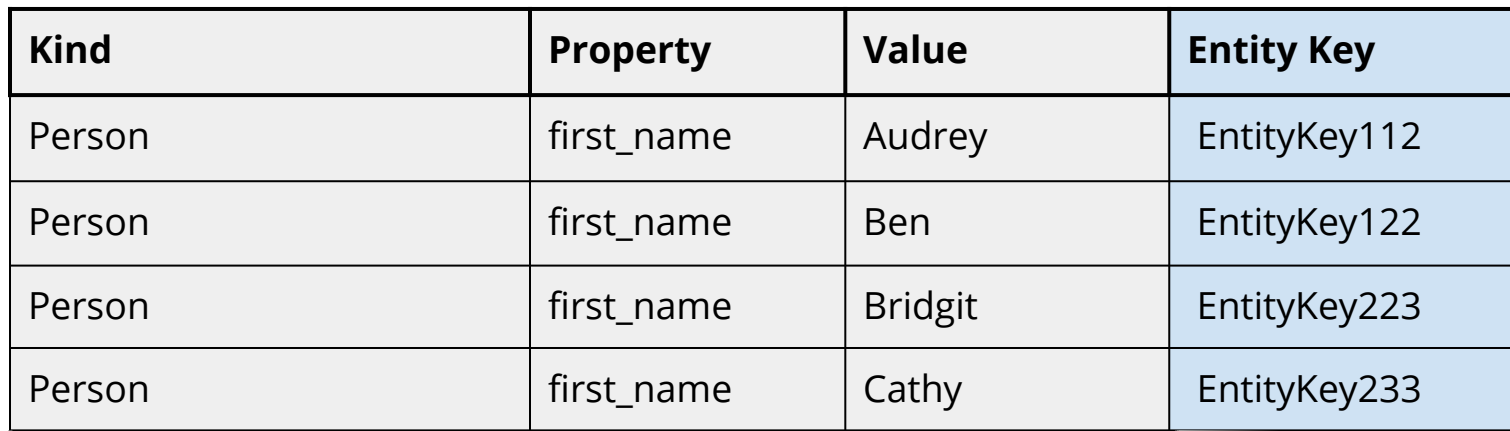
# Index

## Single-property index

- Query for kind: Person
  - `first_name >= "A" and first_name < "C"`
- Scan with prefix [Person first\_name A, Person First\_name B]

**Two** single-property indices are created automatically:

- ascending
- descending



Kind	Property	Value	Entity Key
Person	first_name	Audrey	EntityKey112
Person	first_name	Ben	EntityKey122
Person	first_name	Bridgit	EntityKey223
Person	first_name	Cathy	EntityKey233

**Key of Index table**

**Value**

# Index

Single-property indexes can support queries with:

- Equality filters on one or more properties (merge join)

**first\_name = 'Bob' AND last\_name = 'James'**

- Inequality filters on one property

**first\_name >= 'B' AND first\_name < 'C'**

**first\_name != 'Bob'**

will be executed as

**first\_name < 'Bob' OR first\_name > 'Bob'**

- One sort order

**ORDER BY last\_name ASC**

# Index

How about more  
complex queries?

Use *Composite Indexes*  
(Custom Indexes)

# Index


## Composite index example

- Query for kind : Person

last\_name = "Smith"

first\_name > "A" and first\_name < "D"

Equality filter +  
Inequality filter



<i>Kind /</i>	<i>last_name /</i>	<i>first_name</i>
Person /	Raley /	Jane
Person /	Smith /	Ben
Person /	Smith /	Cathy
Person /	Smith /	Daniel
Person /	Thomas /	Alice

Composite  
Index must  
be **explicitly  
configured**.



# Index

## How to create indexes

- App Engine creates single property indexes for all properties
- You can run queries in the development server to create custom indexes
- You can create or edit index configuration file

# Index

Configuration files for composite index

## Java

- XML

WEB-INF/**datastore-indexes.xml**

WEB-INF/appengine-generated/**datastore-indexes-auto.xml**

- yaml

WEB-INF/**index.yaml**

## Python

- **index.yaml**

Reference: Datastore Indexes ([Java](#)) ([Python](#))

# Index

## Index configuration example

- Query for kind : Person  
last\_name = "Smith"  
first\_name > "A" and first\_name < "D"

<i>Kind</i>	<i>last_name</i>	<i>first_name</i>
Person	Raley	Jane
Person	Smith	Ben
Person	Smith	Cathy
Person	Smith	Daniel
Person	Thomas	Alice

index.yaml

```
- kind: Person
  properties:
    - name: last_name
      direction: asc
    - name: first_name
      direction: asc
```

# Index

## Indexes for multi-valued properties

- An index entry is created for EVERY value of a property

Entity kind Person  
name = Brian  
lucky\_number = {1, 5, 7, 9}

Kind / property / value

Person / lucky\_number / 1

Person / lucky\_number / 5

Person / lucky\_number / 7

Person / lucky\_number / 9

# Index

## Multi-valued properties in Queries

- Multi-valued properties match a query
  - IF AT LEAST ONE value matches ALL the filters

Entity kind Person  
name = Brian  
lucky\_number = {1, 5, 7, 9}

Matches query for?  
Kind is Person  
lucky\_number > 2 and  
lucky\_number < 6

OK

Matches query for?  
Kind is Person  
lucky\_number = 1

OK

Matches query for?  
Kind is Person  
lucky\_number > 1 and  
lucky\_number < 5

X

# Index

Look out for Combinatorial Explosion!

Entity kind Person

name = Brian

number = {1, 5, 7, 9}

colors = {red, blue, green}

friends = {anna, ben, cathy, ...zoe}

name numbers colors friends				
Brian	9	red	anna	
Brian	9	red	ben	
Brian	9	red	cathy	
...	...	...	...	
Brian	9	red	zoe	

name numbers colors friends				
Brian	9	blue	anna	
Brian	9	blue	ben	
Brian	9	blue	cathy	
name numbers colors friends				
Brian	9	green	anna	
Brian	9	green	ben	
Brian	...			
...	...			
Brian	9	green	zoe	

$4 \times 3 \times 26 = 312$  index entries



# Restrictions and Workarounds for Queries

# Query

Sounds like Datastore Queries are not as flexible as RDBMS queries.

True. It's very important to understand the **restrictions** of Datastore Query and workarounds for them *before* you start designing your data model.



# Query: Missing Properties

- Entities with **no property or an** unindexed value are **not included in results**

Query for:

Kind = Person

last\_name != Arundel

**Missing Property**  
is not equal to **Null/None**



Kind	last_name	first_name
Person	Anderson	Jane
Person	Artwood	
Person		Jenny

# Query: Inequality Filters

- Inequalities filters: limited to one property per query

Query for:

first\_name = Cathy

last\_name > Able

last\_name < Mooney

OK

Query for:

first\_name > Cathy

last\_name > Able

X

# Query: Inequality Filters and Sorting

- A property with an inequality filter must be **sorted first**

Query for:

first\_name = Cathy

last\_name > Able

sort by last\_name

OK

Query for:

last\_name > Able

sort by first\_name

X

# Query: Restrictions on Inequality Filters


**Why** the restrictions on inequality filters ?

- To avoid having to scan the entire index table, the query mechanism relies on all of a query's potential results being able to be converted to a key range scan on the index

Just suppose we allowed...

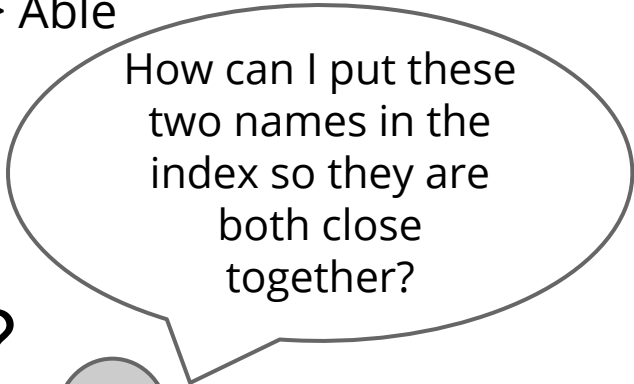
first\_name > Aaron

last\_name > Able



Why are inequality filters so restrictive?

Zoe Acton  
Ben Yodel




How can I put these two names in the index so they are both close together?



# Query: Restrictions on Inequality Filters

**Why** the restrictions on inequality filters ?

- To avoid having to scan the entire index table, the query mechanism relies on all of a query's potential results being able to be converted to a key range scan on the index

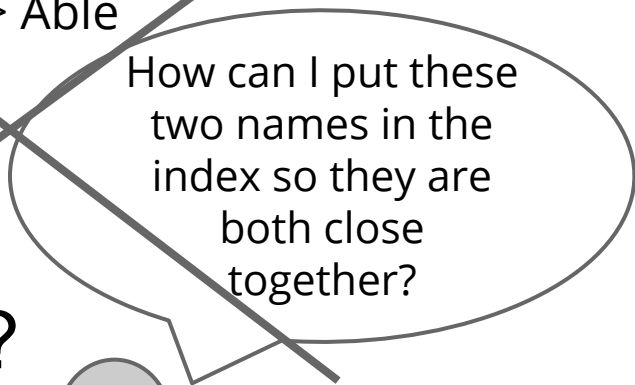


Why are inequality filters so restrictive?

~~Just suppose we allowed...~~

~~first\_name > Aaron  
last\_name > Able~~

Zoe Acton  
Ben Yodel




How can I put these two names in the index so they are both close together?



# Query: JOIN Operation

- JOIN operation is not supported
- Use **Denormalization**  
It's a known practice for any scalable database design

```
SELECT * FROM PERSON p, ADDRESS a  
WHERE a.person_id = p.id AND p.age > 25  
AND a.country = "US"
```



```
SELECT FROM Person  
WHERE age > 25 and country = "US"
```



Maintain **country** in Person

# Query: Aggregation

- Aggregation queries are not supported
- Datastore does not support aggregation queries (group by, having, sum, avg, max, min)
- Use a special entity that maintains aggregated values
  - Counter entity
  - Be careful not to make the entity bottleneck (by 1 updates/sec limit)
  - Use **Sharding Counter** pattern
  - Or **Memcache putIfUntouched() + Datastore insert**

# Query: Aggregation (continued)

- Use batch processing to aggregate values asynchronously
  - Backend instance
  - [App Engine MapReduce](#)
- [Datastore Statistics](#)
  - For counting entities
  - Updated once per day
- Use Sorting for MIN() or MAX()
  - Sort by a property: the first entity will have min/max value.





# Cost of Indexing

# Cost of Index

- **Index** consumes **Datastore space & instance hours**
- Take the **cost of Index** into account for **cost estimation**  
Read:
  - [Understanding Write Costs](#)
  - [How Entities and Indexes are Stored](#)
- New Index for a large set of entities may take a **long time**

# Datastore Statistics

- View statistics in the Admin Console
  - Datastore > Datastore Statistics

Display statistics for:  
Kind: **All Entities**

Statistics are updated at least once per day. [Learn more](#)

Last updated: 3 days, 15:25:16 ago

<b>Entities</b>	<b>Built-in Indexes</b>	<b>Composite Indexes</b>
25 MBytes	169 MBytes	44 MBytes
91,634	1,461,907	218,428

# Index : Cost

Setting a property without indexing

## Python

```
class Conference (db.Model):  
    main_contact = db.StringProperty(indexed=False)
```

## Java

```
conference.setUnindexedProperty("mainContact", "Adam Bolivar");
```

- Don't index long strings (use Search API)
- Text and Blob are automatically treated as unindexed

# Index : Cost

## Deleting Indexes

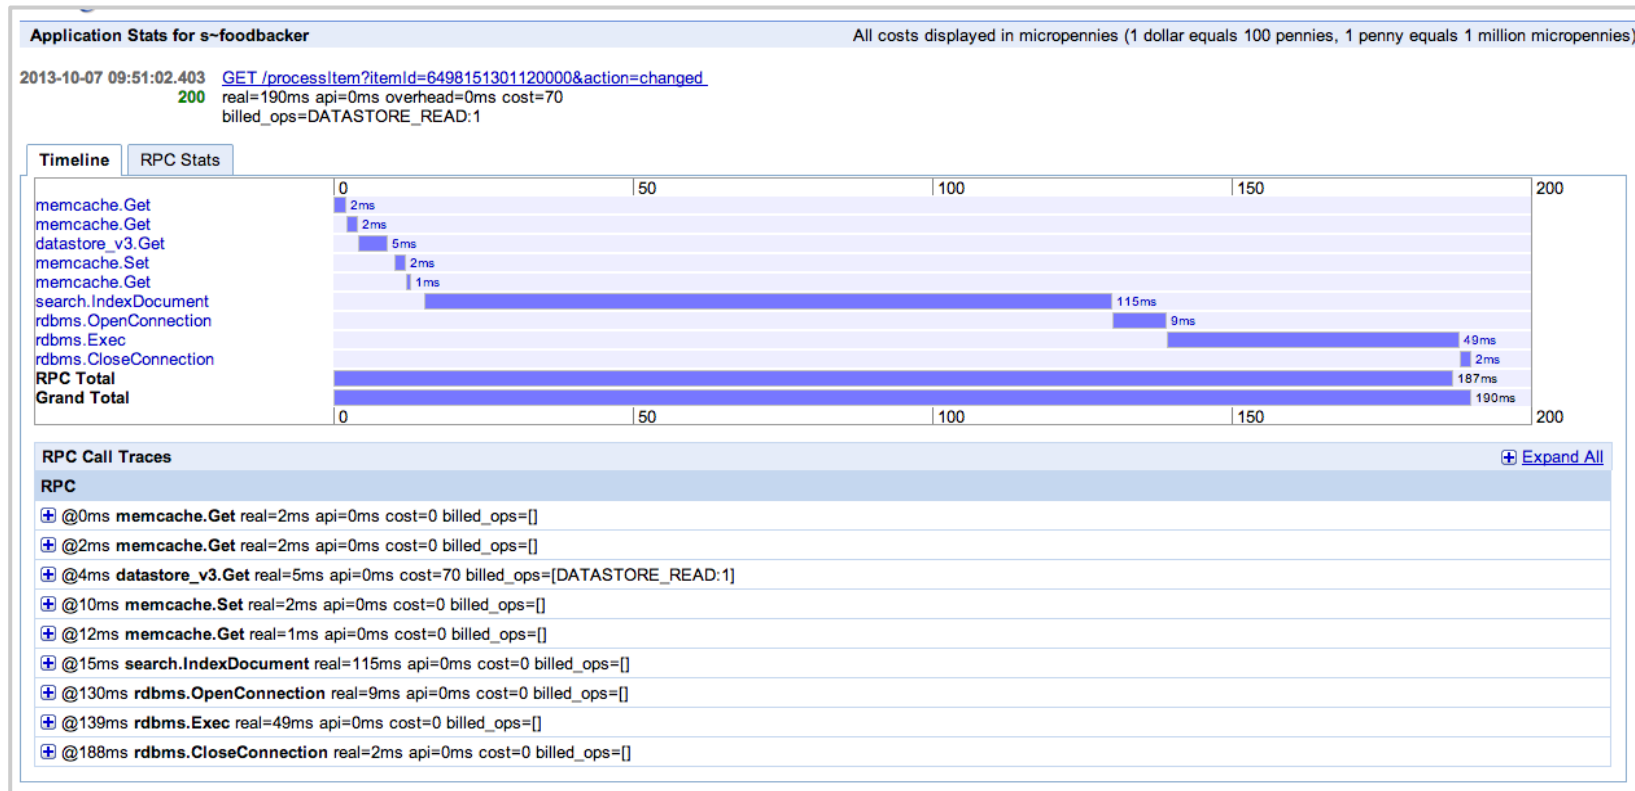
- Existing indexes remain when index config changes
- Lets you:
  - Leave an older version of the app running while new indexes are being built
  - Revert to the older version if needed

To delete all unused indexes:

- Update index config file
- Run:

```
appcfg.sh vaccum_indexes myapp
```

# Profiling with AppStats Tool



# Enable AppStats Recording

## Java

web.xml

```
<filter>
  <filter-name>appstats</filter-name>
  <filter-class>com.google.appengine.tools.appstats.AppstatsFilter</filter-class>
</filter>
<filter-mapping>
  <filter-name>appstats</filter-name>
  <url-pattern>/*</url-pattern>
</filter-mapping>
```

## Python

appengine\_config.py

```
def webapp_add_wsgi_middleware(app):
    from google.appengine.ext.appstats import recording
    app = recording.appstats_wsgi_middleware(app)
    return app
```

# Enable AppStats Admin Console Page

## Java

web.xml

```
<servlet>
  <servlet-name>appstats</servlet-name>
  <servlet-class>com.google.appengine.tools.appstats.AppstatsServlet</servlet-class>
</servlet>
<servlet-mapping>
  <servlet-name>appstats</servlet-name>
  <url-pattern>/appstats/*</url-pattern>
</servlet-mapping>
```

## Python

app.py

```
builtins:
- appstats: on
```




# Quiz

Which of the following methods can be used to create datastore indexes?  
(select **3** of the available options)

- ☐ App Engine automatically creates single property indexes for all properties
- ☐ You need to manually run queries in the development server to create single property indexes
- ☐ You can run queries in the development server to create custom indexes
- ☐ App Engine automatically creates custom indexes when deploying your application
- ☐ You can edit index configuration files manually
- ☐ None of the above

# Codelab

- Try out multiple queries in the same source code as the previous Hands on exercise



Images by Connie Zhou

# Resources

- **Developer Documentation**

- Datastore Queries ([Java](#)) ([Python](#))
- Transactions ([Java](#)) ([Python](#))

- **App Engine Datastore Articles**

- Mastering the Datastore series
- [Transaction Isolation](#)
- [Handling Datastore Errors](#)
- Understanding Write Costs
- [How Entities and Indexes are Stored](#)

# Resources

- **Google I/O Sessions**

- [Under the Covers of the Google App Engine Datastore, Ryan Barret, Google I/O 2008](#)
- [Building Scalable, complex Apps on App Engine, Brett Slatkin, Google I/O 2009](#)
- [Next gen queries, Alfred Fuller, Google I/O 2010](#)





cloud.google.com