



# Decoupling Work Using Queues & Scheduled Tasks

Google App Engine

# Agenda

1

Task Queue Overview

2

Push Queue

3

Pull Queue

4

Cron or Scheduled Tasks


5

Lab Exercise

# Task Queues

- Task is a unit of work to be executed by your application or external workers.
  - Executed asynchronously outside a user request
  - Use Cases:
    - Sending e-mail
    - Batching datastore writes
- Task Queue is an App Engine service to manage and run tasks.

# Task Queues



# Two Kinds of Task Queue Services

## Push queues

- App Engine managed service
  - App Engine schedule/execute/delete/retry automatically
  - Your application creates tasks with HTTP request handlers
- Less control and only directly accessible from your App Engine app

## Pull queues

- Manual service
  - Your application schedules, executes, deletes, and retries task
- More control and available from outside App Engine app

# How to Configure Queues?

## Configure queues in a queue config file

- Java: `queue.xml`
- Python: `queue.yaml`

## In the queue config file, you can specify

- Total storage limit (Optional safety check)
- For each queue,
  - Queue name (Required)
  - Mode ( Push/Pull)
  - Other parameters specific to each queue type

config docs

- [Java](#)
- [Python](#)

# Queue Configuration Files

## Java

```
<queue-entries>
  <queue>
    <name>push-queue-1</name>
    <rate>1/s</rate>
  </queue>
  <queue>
    <name>pull-queue-1</name>
    <mode>pull</mode>
  </queue>
</queue-entries>
```

queue.xml

## Python

```
queue:
- name: push-queue-1
  rate: 1/s
- name: pull-queue-1
  mode: pull
```

queue.yaml

# Adding A Task to A Queue

1. Get queue
2. Create task with options
3. Add task to queue

## Java

```
Queue queue = QueueFactory.getQueue("Qname");  
TaskOptions task = TaskOptions.Builder(...)  
queue.add(task);
```

## Python

```
queue = taskqueue.Queue("Qname")  
task = taskqueue.Task(...)  
queue.add(task)
```



# Adding Tasks in A Datastore Transaction

- Tasks can be enqueued during a transaction
- Task does not *execute* within the transaction

## Java

```
Transaction txn = ds.beginTransaction();  
queue.add(TaskOptions.Builder.withUrl("/handler"));  
txn.commit();
```

## Python

```
def a_method(parameter...):  
    taskqueue.add(url='/handler', transactional=True)  
  
db.run_in_transaction(a_method, parameters...)
```



# Push Queues

# Push Queue Overview

- Task Definition
  - Name, parameters, headers, and payload etc. with max 100KB task size
  - Specify handler (internal URLs) and method (GET/POST)
- Task Scheduling Per Queue
  - Executed ASAP and as many as possible with *token bucket algorithm*
  - Processed in FIFO order (mostly), ETA or countdown
- Task Execution:
  - Tasks processed as HTTP requests (GET/POST)
  - May cause new instances to be created
  - 10 minute limit for request processing for automatic scaling
  - 24 hour limit for request processing for basic and manual scaling
  - Task is deleted after successful execution, back to queue if fails

# Push Queue Task Definition

## Specify options in code when creating tasks

- **url**
- task name
- payload
- params
- headers
- countdown / eta
- target
- retryOptions

- TaskOptions [JavaDoc](#)
- Task Reference [Python](#) (*includes nice summary of push vs pull options*)

# Order of Task Execution

Order of task execution is **best effort** based on

- Position of task in the queue (FIFO)
- Task's [eta](#) or [countdown](#) property
  - Mutually exclusive
  - Task is executed on or after eta or countdown
- Backlog of tasks in the queue
  - New task might queue jump to optimize execution in case of big backlog

# Default Queue

- App Engine has a default push queue
- No need to configure it—uses default settings
- If you want to configure it:
  - Define queue named default in queue config file

## Java

```
Queue queue = QueueFactory.getDefaultQueue();  
Queue queue = QueueFactory.getQueue("Qname");
```

## Python

```
queue = taskqueue.Queue()  
queue = taskqueue.Queue("Qname")
```

# Tuning Task Scheduling

## Parameters for token bucket algorithm:

- rate -- usual rate
- bucket-size -- cap on peak demand
- max concurrent requests

## Example:

- rate = 10
- bucket-size = 40

Queue can execute 10 tasks per second. If rate is not used up, accrue up to max of 40 tasks.

# What Happens if a Task Fails?

- HTTP 200 - 299 response code for success; otherwise fails.
- If a task fails, it goes back in the queue
- App Engine retries it until it succeeds
- Specify task retry options
  - In queue config file
  - In code when creating the task



# Task Retry Options

- Retry Limits
  - **task-retry-limit** - min retries
  - **task-age-limit** - min elapsed time to keep retrying
    - If both task-retry - limit and task-age-limit are reached, task is deleted
- Backoff Strategy
  - **min-backoff-seconds** -- min amount of time to wait before the first retry
  - **max-backoff-seconds** -- max delay between retries
  - **max-doublings** -- max times to double delay

# Push Queue in Development Server

- The dev server executes tasks at the appropriate time
- Executes tasks as close to their ETA as possible

## ***However***

- Does not respect rate and bucket-size
- Setting a rate of 0 does not prevent tasks from being executed automatically
- Does not retry tasks
- Does not preserve queue state across server restarts

# Push Task Request Headers

- X-AppEngine-QueueName
- X-AppEngine-TaskName
- X-AppEngine-TaskRetryCount
- X-AppEngine-ExecutionCount
- X-AppEngine-TaskETA
- X-AppEngine-FailFast: failing without spawning new instances

***These headers are set internally by App Engine. Stripped out if request comes from outside App Engine.***

# Push Queue Best Practices

- Set security so only admins can access handler URLs
- Use different queues for different purposes
  - Allows rates to be fine-tuned for performance/cost trade-off
- Batch task creation (< 100 tasks per batch)
- If order is critical, implement a control mechanism
- Specify target to use a specific version
  - Unless specified otherwise, task executes on the version of the app that spawned it



## Pull Queues

# Pull Queue Overview

- Task Definition (< 1MB)
  - Do not specify handler, target, method (always PULL)
  - Tasks can have tags; you can lease tasks by tag
- Task Queue Definition
  - Set mode to PULL
  - No default pull queue
- Task Execution
  - Task leased by worker - no scheduling, no auto retry
  - Workers must delete tasks -- no auto deletion
- REST interface (with ACL)

# Pull Queue Task Options

- task name
- **method (PULL)**
- payload
- params
- tags (**Beta**)
- task retry limit

- TaskOptions [JavaDoc](#)
- Task Reference [Python](#) (*includes nice summary of push vs pull options*)

# Processing Pull Tasks

- Lease Tasks
  - Workers lease tasks from the queue for a period of time
  - Lease up to 1000 tasks at a time
  - Modify lease to extend duration
- Delete Tasks
  - Worker must explicitly delete task when done
  - If task is not deleted when lease expires, it goes back in the queue
- Retry Tasks = Lease Tasks Again



# Pull Queue Best Practices

- Scale up and down workers yourself
- Don't create bloated tasks
  - Limit of 32MB of task data per lease request
- Distribute task execution time, if possible.
- Choose a lease close to the worst-case scenario
- Catch transient and deadline exceeded exceptions and backoff

# RESTful APIs for Pull Queue

## Taskqueues

- [get](#) -- Get detailed information about a TaskQueue

## Tasks

- [delete](#) -- Deletes a task from a TaskQueue
- [get](#) -- Gets the named task in a TaskQueue
- [insert](#) -- Insert a task into an existing queue
- [lease](#) -- Acquires a lease on the topmost N unowned tasks in a queue
- [list](#) -- Lists all non-deleted Tasks in a TaskQueue, whether or not they are currently leased, up to a maximum of 100
- [patch](#) -- Update tasks that are leased out of a TaskQueue
- [update](#) -- Update the duration of a task lease



# Google App Engine

## Cron And Scheduled Tasks

# Cron Tasks

- Use cron for regularly scheduled tasks
- Cron tasks use a queue named "\_\_cron"
- Cron job is executed as a GET request to the specified URL
- Configure in cron.xml or cron.yaml
  - URL
  - Schedule

# Cron Configuration

- Url
  - The url to call (escape &, <, >, ', ")
- Schedule
  - The times/dates to execute the task
  - From/to or **synchronized** task
- Timezone
  - Optional, the standard zoneinfo name (defaults to UTC)
- Target
  - Optional, the target version of application (defaults to the default)

# Cron Configuration Samples

## Java cron.xml

```
<cronentries>
  <cron>
    <url>/recache</url>
    <description>
      Repopulate the cache
    </description>
    <schedule>
      every 2 minutes
    </schedule>
  </cron>
</cronentries>
```

## Python cron.yaml

```
cron:
- description: new hourly summary
  job
    url: /tasks/summary
    schedule: every 1 hours
```



**url** and **schedule** are required parameters

# Cron Tips

- Set security so only admins can access task handler URLs
- Go to the cron URL in a browser to test it
- When a cron task is executed, it has the header "X-AppEngine-Cron: true".
- You can have
  - 20 cron jobs for free apps
  - 100 for paid apps
- Review cron jobs in the Admin Console

# Quiz

If you wanted to share some work from App Engine with **Compute Engine** how could go about it using Task Queues?


(pick **one** answer)

- ☐ A push queue
- ☐ Use Cron
- ☐ A scheduled task
- ☐ A pull queue
- ☐ All of the above
- ☐ It's not possible to share work in this way



# Codelab

- Use a push queue to simulate sending out mail for a newly created conference
- When a conference is created, add a task to the task queue to notify all users about the new conference
- Write the handler to notify the appropriate users



Images by Connie Zhou

# Review Tasks in Admin Console

## Push Queues

Queue Name	Max Rate	Enforced Rate	Bucket Size	Tasks in Queue	Run Last Min	Running
<a href="#">default</a>	1.0/s	1.00/s	5.0	0	1	0
<a href="#">notify-queue</a>	1.0/s	1.00/s	5.0	0	1	0

## Pull Queues

Queue Name	Oldest Task	Tasks in Queue	Leased in Last Min
<a href="#">review-conference-queue</a>	2014-01-10 15:00:30 (0:00:00 ago)	3	3

Queue Name	Maximum Rate	Enforced Rate	Bucket Size	Tasks in Queue	Run in Last Min	Running
notify-queue	1.0/s	1.00/s	5.0	10	15	1

Purge Queue

Delete Queue

Pause Queue

# Resources

- Configuring Queues [\(Java\)](#) [\(Python\)](#)
- Using Push Queue [\(Java\)](#) [\(Python\)](#)
- Using Pull Queues [\(Java\)](#) [\(Python\)](#)
- Article: [Getting Started with Task Queues](#)
- Backend Instances [\(Java\)](#) [\(Python\)](#)
- Scheduling tasks with cron [\(Java\)](#) [\(Python\)](#)





[cloud.google.com](https://cloud.google.com)